SENTIMENT ANALYSIS IN MOVIE COMMENTS

HELLO!

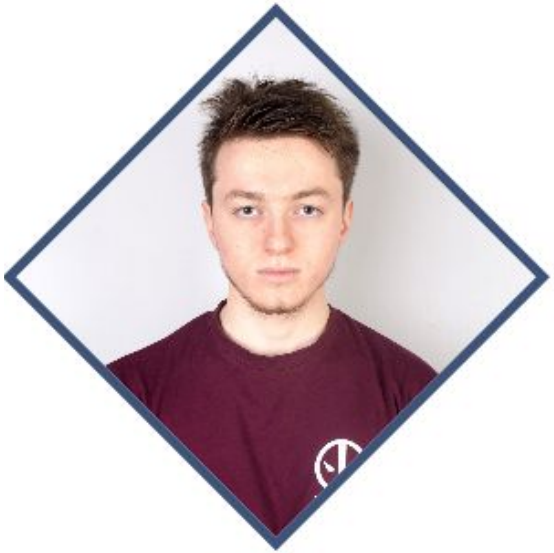

Denis Castéran

Hugo Rybinski

SCIA 2018

GOAL

- Fetch movie datas (especially reviews)
- Analyze the datas to determine if it is a positive or a negative review
- Extract informations from the datas and display it

SOFTWARES & TECHNOLOGIES USED

Languages:

- Scala
- Python

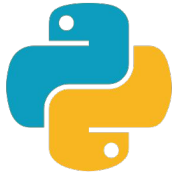

From the Hadoop stack:

- Kafka
- Spark

Misc:

- Spark-notebook
- Docker

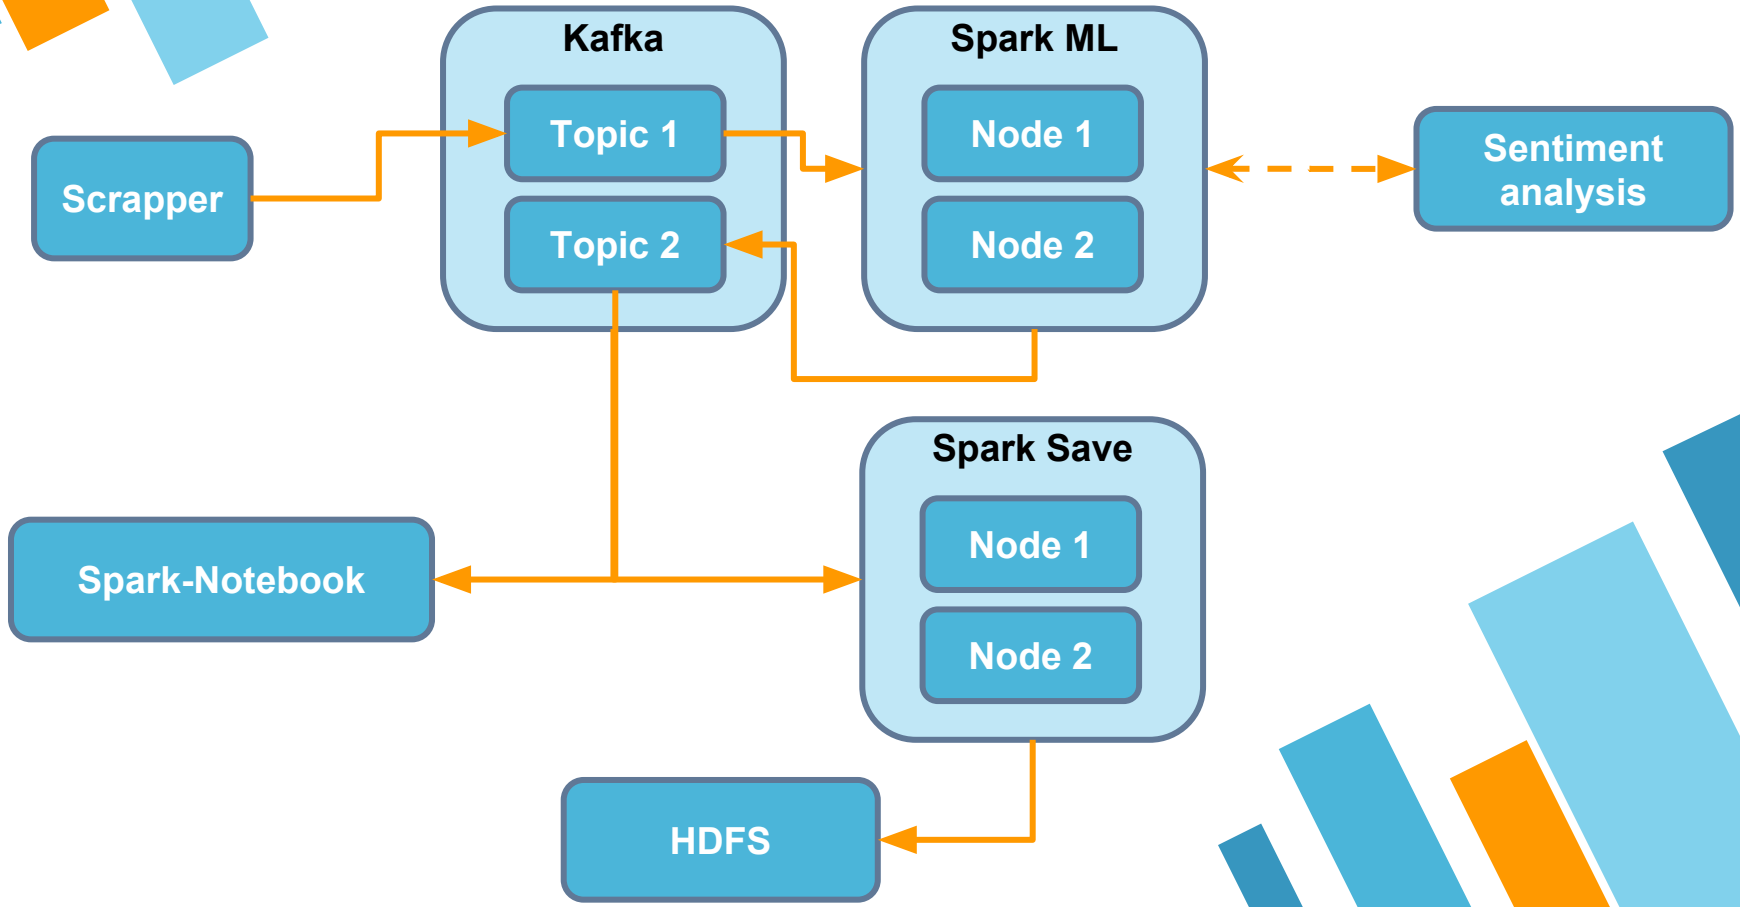
Kafka
Topic 1
Topic 2
Spark ML
Node 1
Node 2
Scrapper
Sentiment analysis
Spark Save
Node 1
Node 2
Spark-Notebook
HDFS

1.

DATABASE SCRAPING & KAFKA

Get the datas

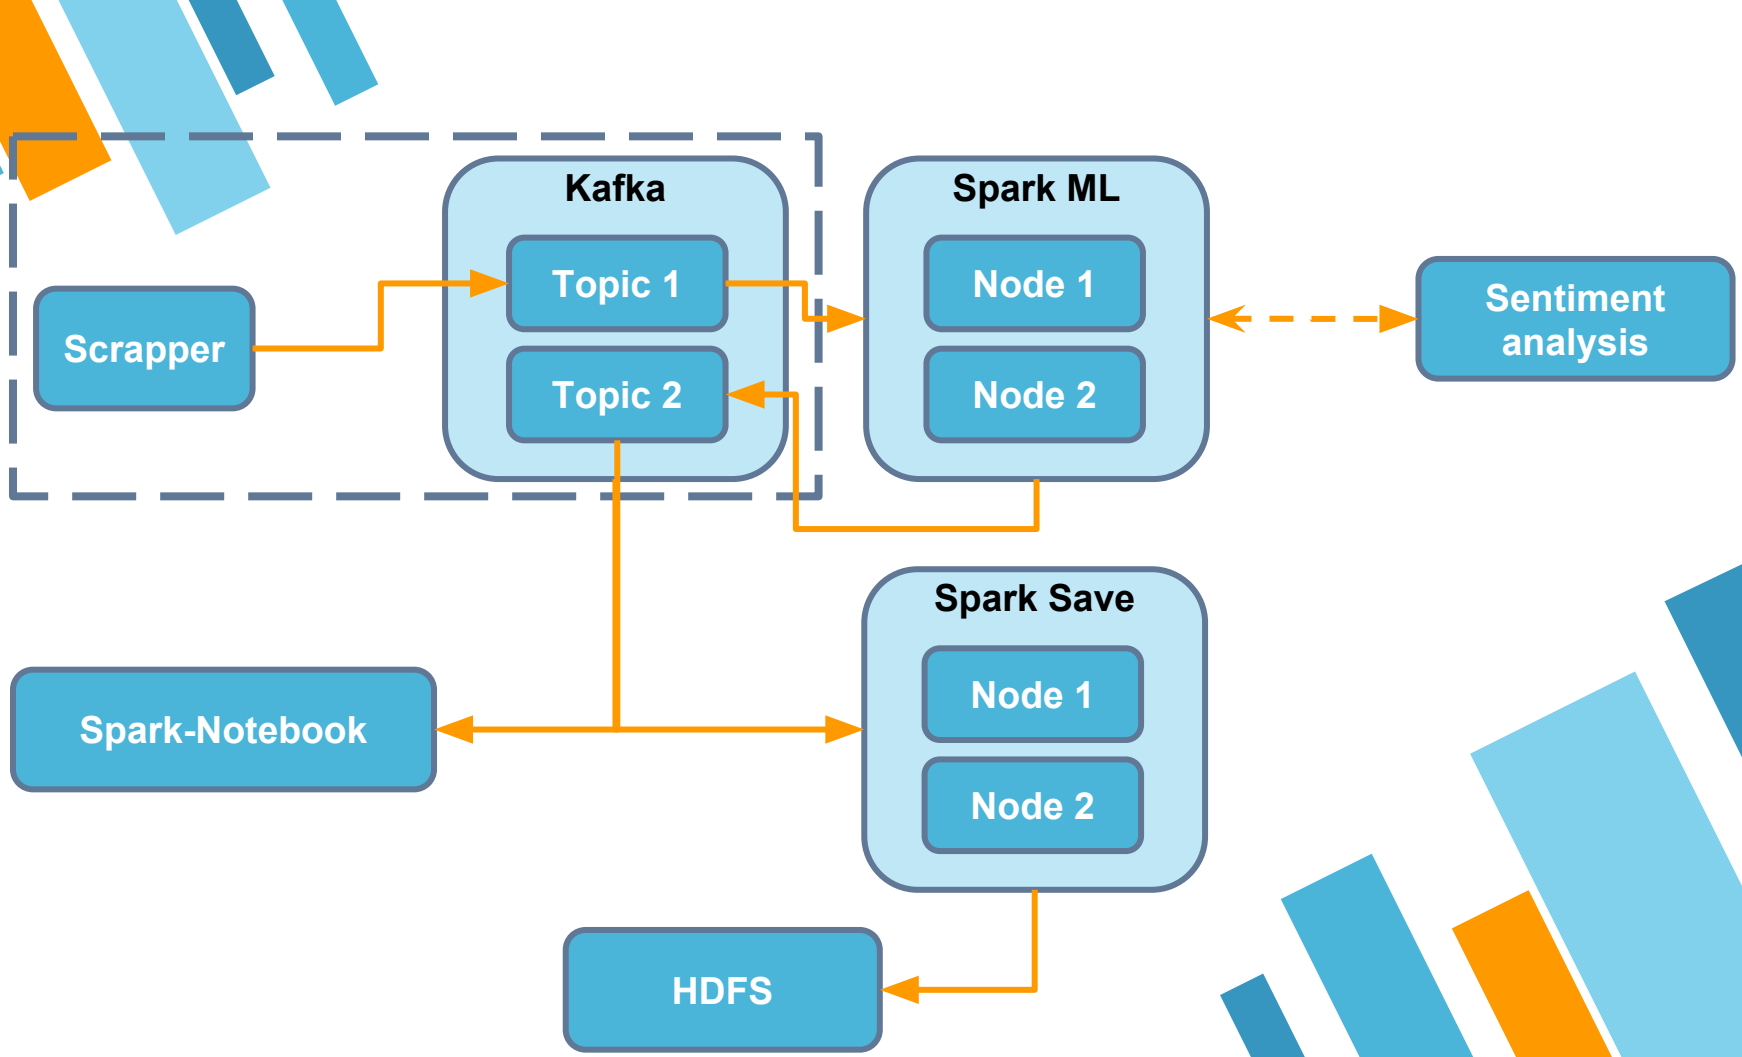
Kafka
Topic 1
Topic 2
Scrapper
Spark ML
Node 1
Node 2
Sentiment analysis
Spark Save
Node 1
Node 2
Spark-Notebook
HDFS

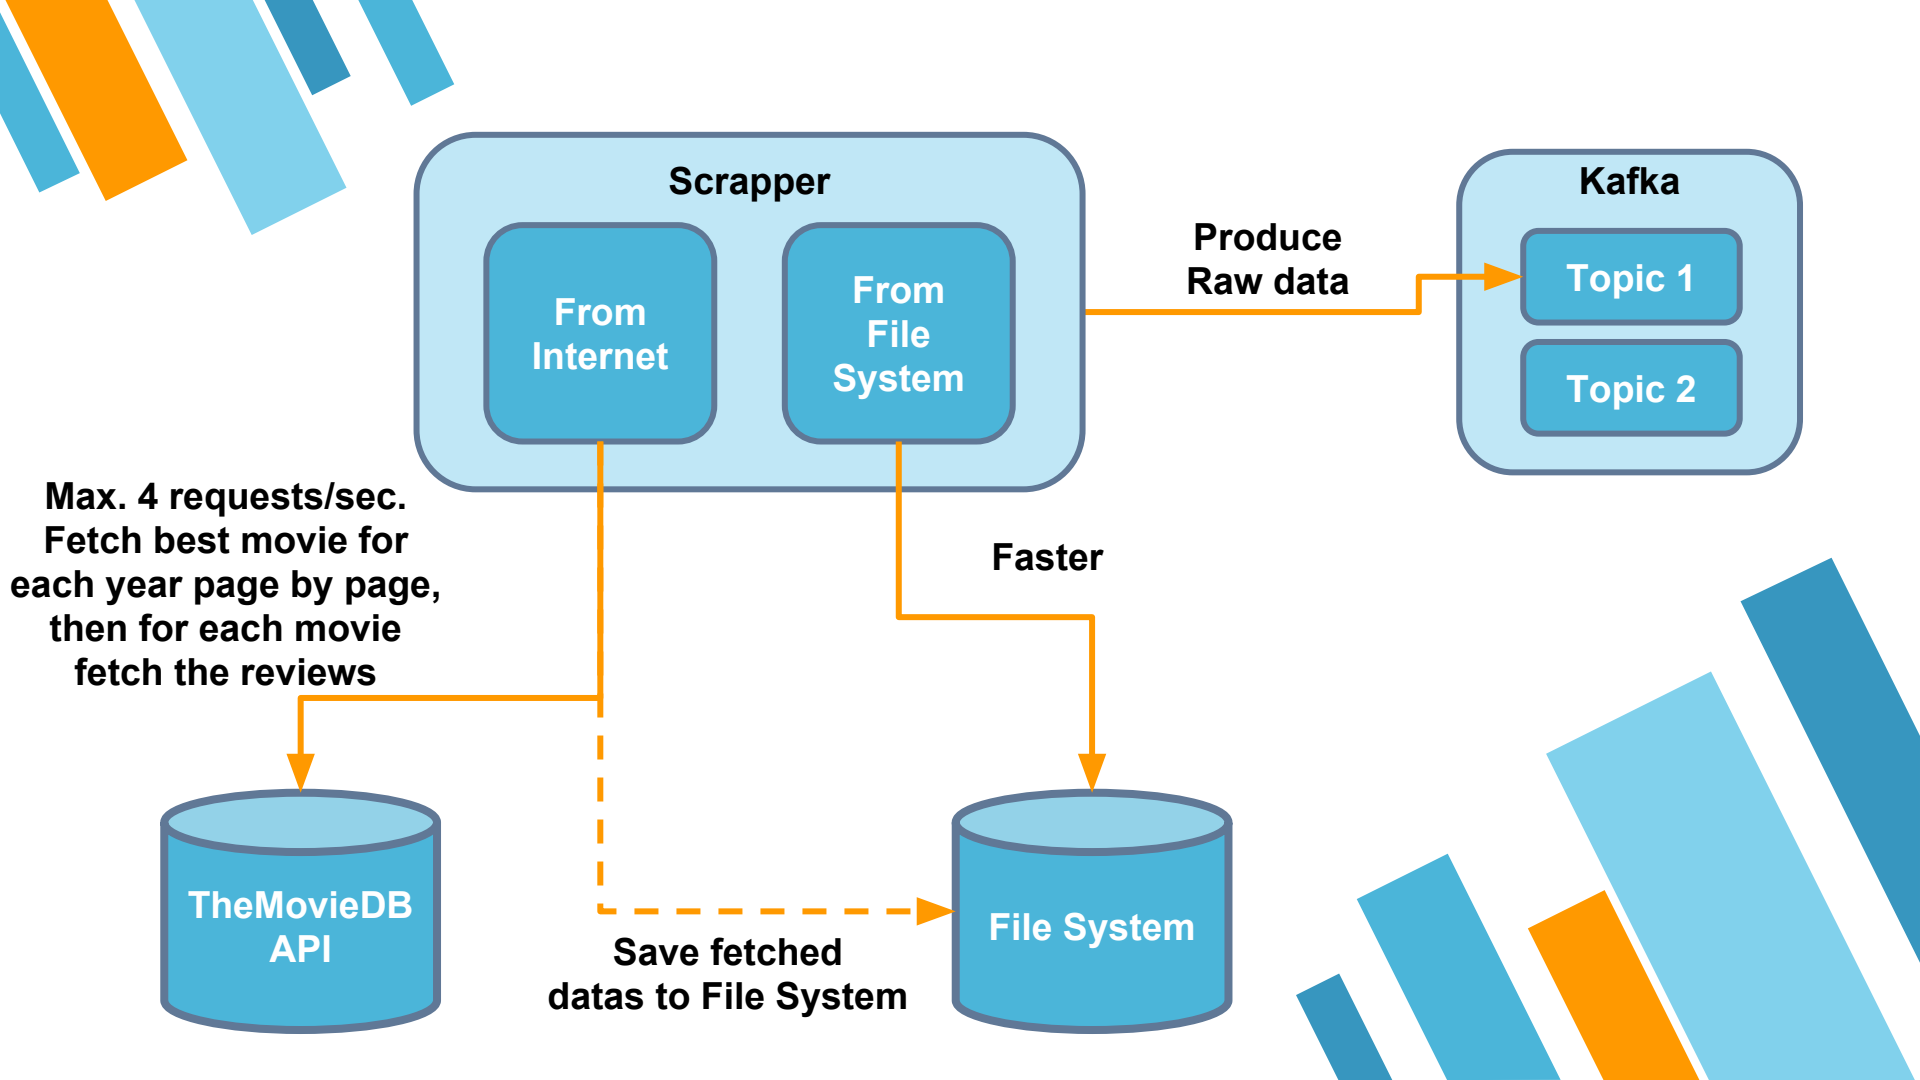
Scrapper
From Internet
From File System
Produce Raw data
Kafka
Topic 1
Topic 2
Max. 4 requests/sec. Fetch best movie for each year page by page, then for each movie fetch the reviews
Faster
TheMovieDB API
Save fetched datas to File System
File System

KAFKA SETUP

- Simple setup with two Docker images:
  - Zookeper
  - Kafka
- Possibility to add broker on other machines:
  - Just start the Kafka image somewhere else, but change the address and port of the Zookeeper server
- Easy to manage with kafka-manager

2.

SPARK TREATMENT

Parallelize machine learning

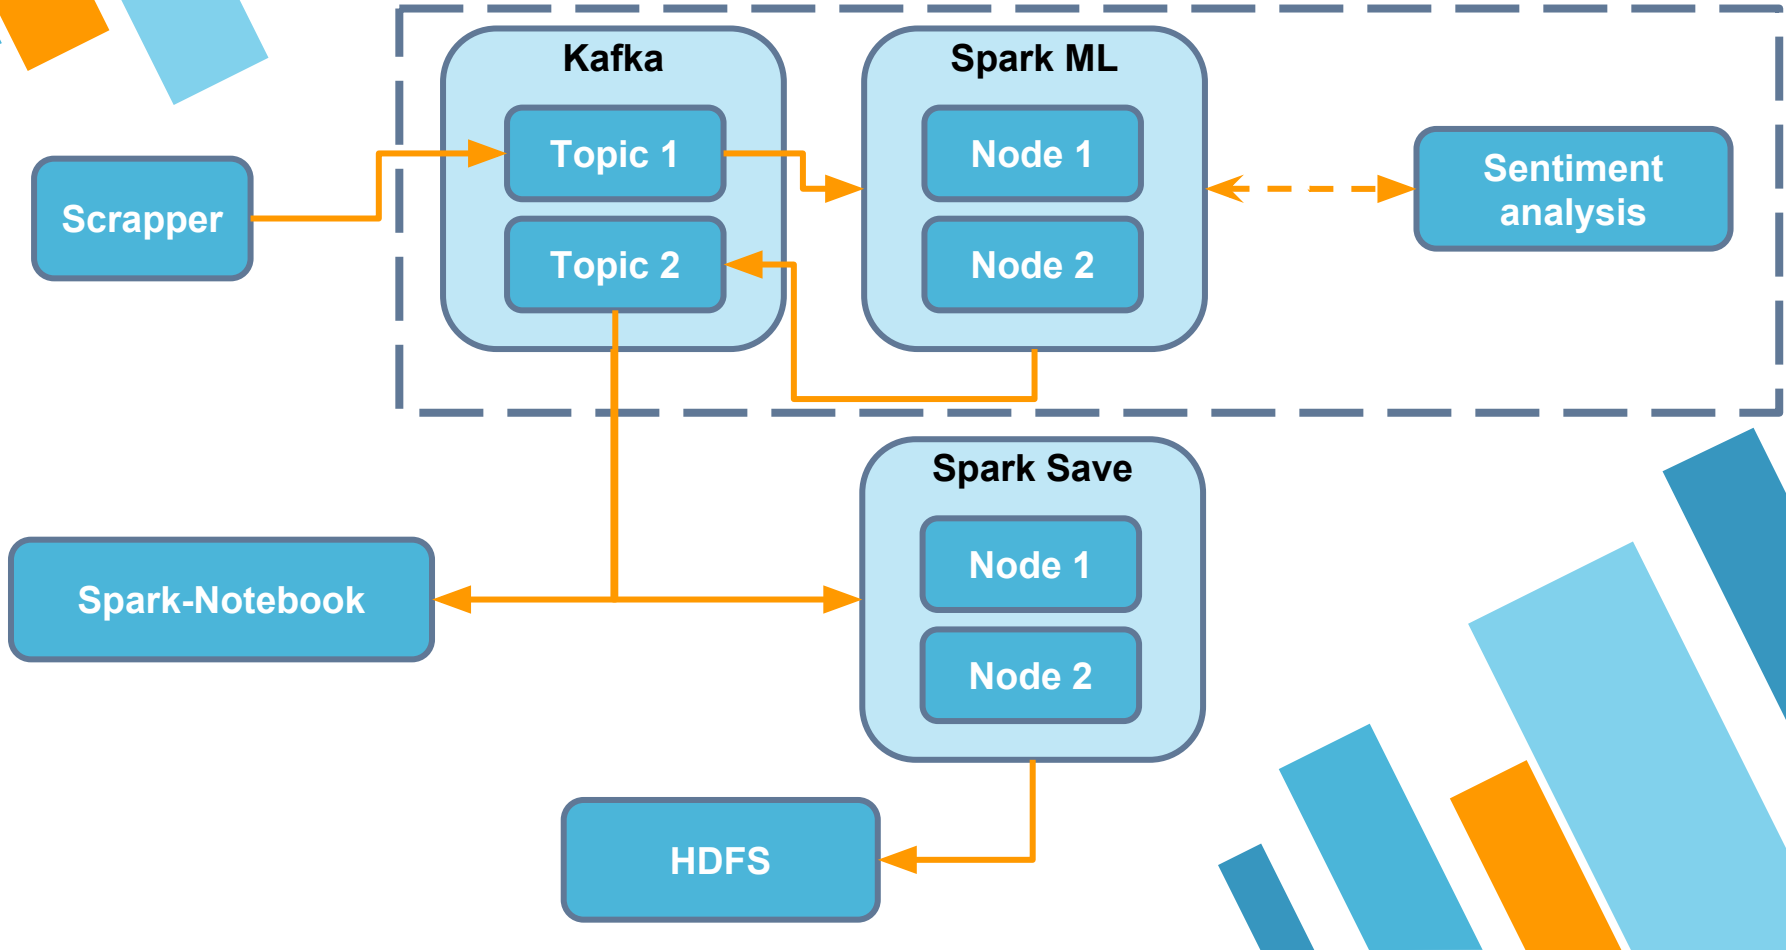
Kafka
Topic 1
Topic 2
Spark ML
Node 1
Node 2
Scrapper
Sentiment analysis
Spark Save
Node 1
Node 2
Spark-Notebook
HDFS

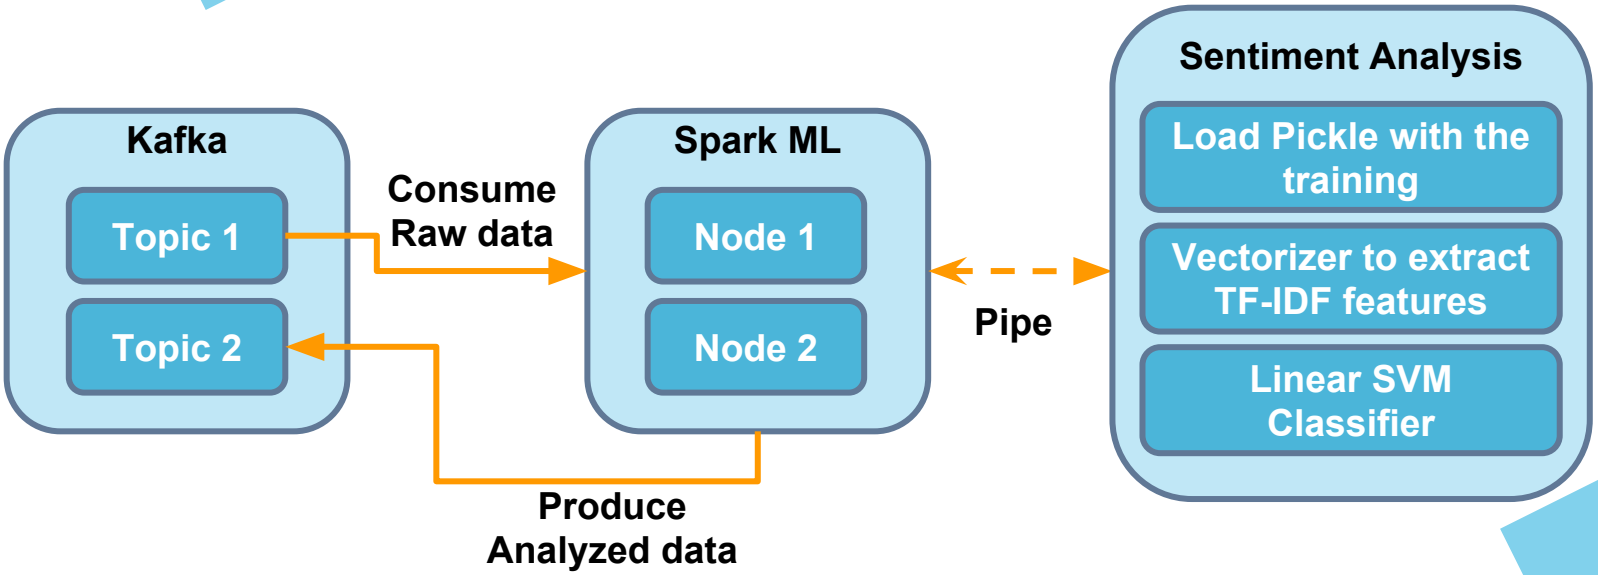
Kafka
Topic 1
Topic 2
Consume
Raw data
Spark ML
Node 1
Node 2
Pipe
Sentiment Analysis
Load Pickle with the training
Vectorizer to extract TF-IDF features
Linear SVM Classifier
Produce
Analyzed data

SPARK

- Spark streaming :
  - Allow parallelized computation
  - Quick computation of sentiment analysis
  - Easy to communicate with kafka

SENTIMENT ANALYSIS

- Built using Scikit-Learn
- Train on local machine once, trained classifier is saved in a Pickle
  - Training dataset : 10 000 labelled reviews from IMDB
  - Testing dataset : 10 000 labelled reviews from IMDB
  - ~80% accuracy

3.

DATA PERSISTENCE

Save the computed datas

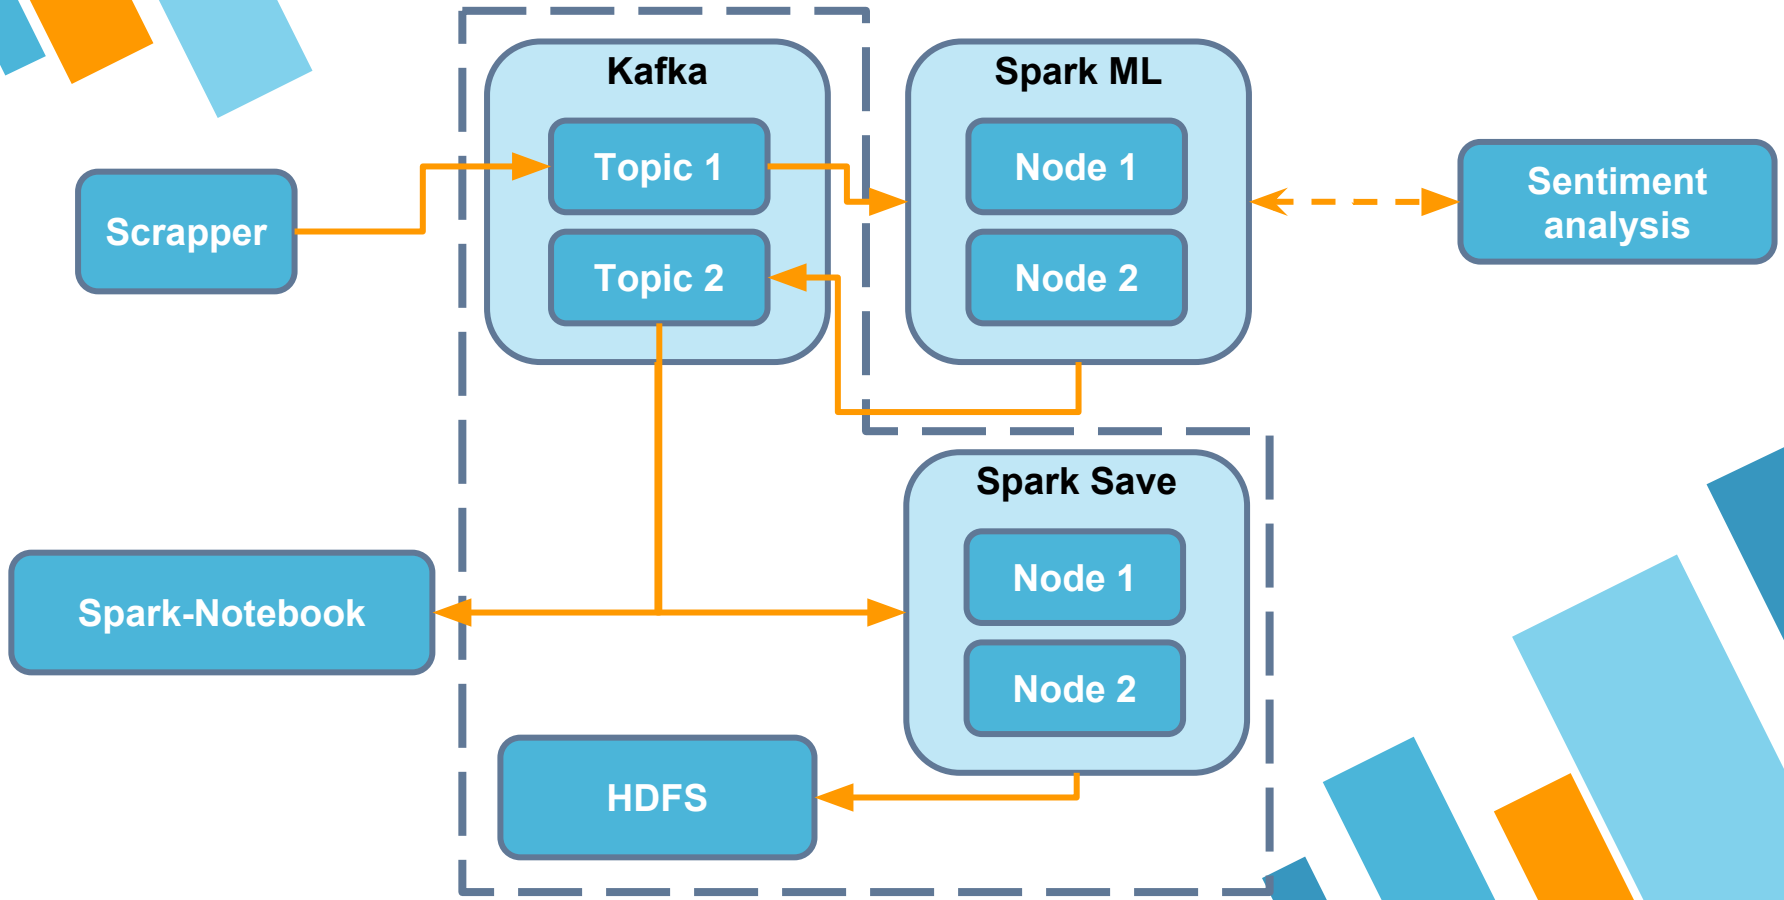
Kafka
Topic 1
Topic 2
Spark ML
Node 1
Node 2
Scrapper
Sentiment analysis
Spark Save
Node 1
Node 2
Spark-Notebook
HDFS

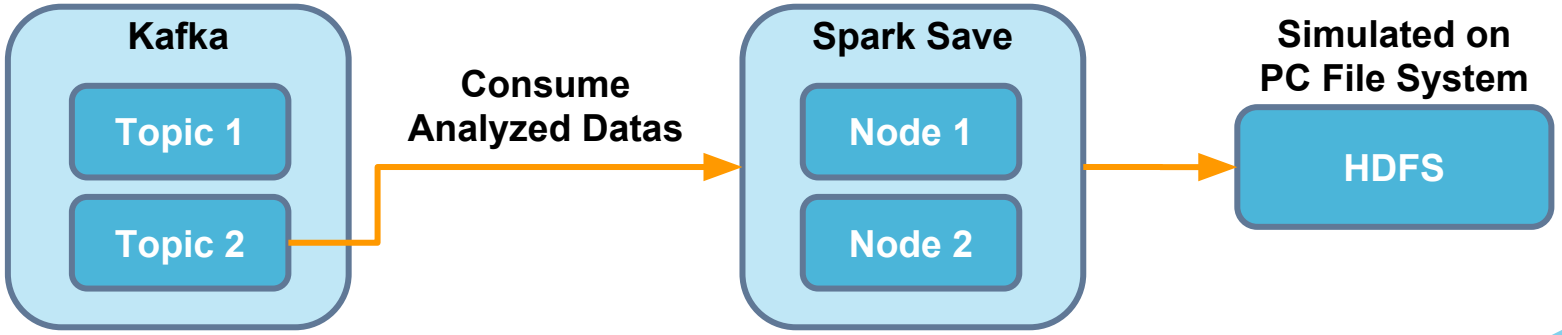
Kafka
Topic 1
Topic 2
Consume
Analyzed Datas
Spark Save
Node 1
Node 2
Simulated on
PC File System
HDFS

HDFS

- Simulated on file system :
  - The dataframe that contains the RDD is saved

4.

DATA DISPLAY

With spark-notebook

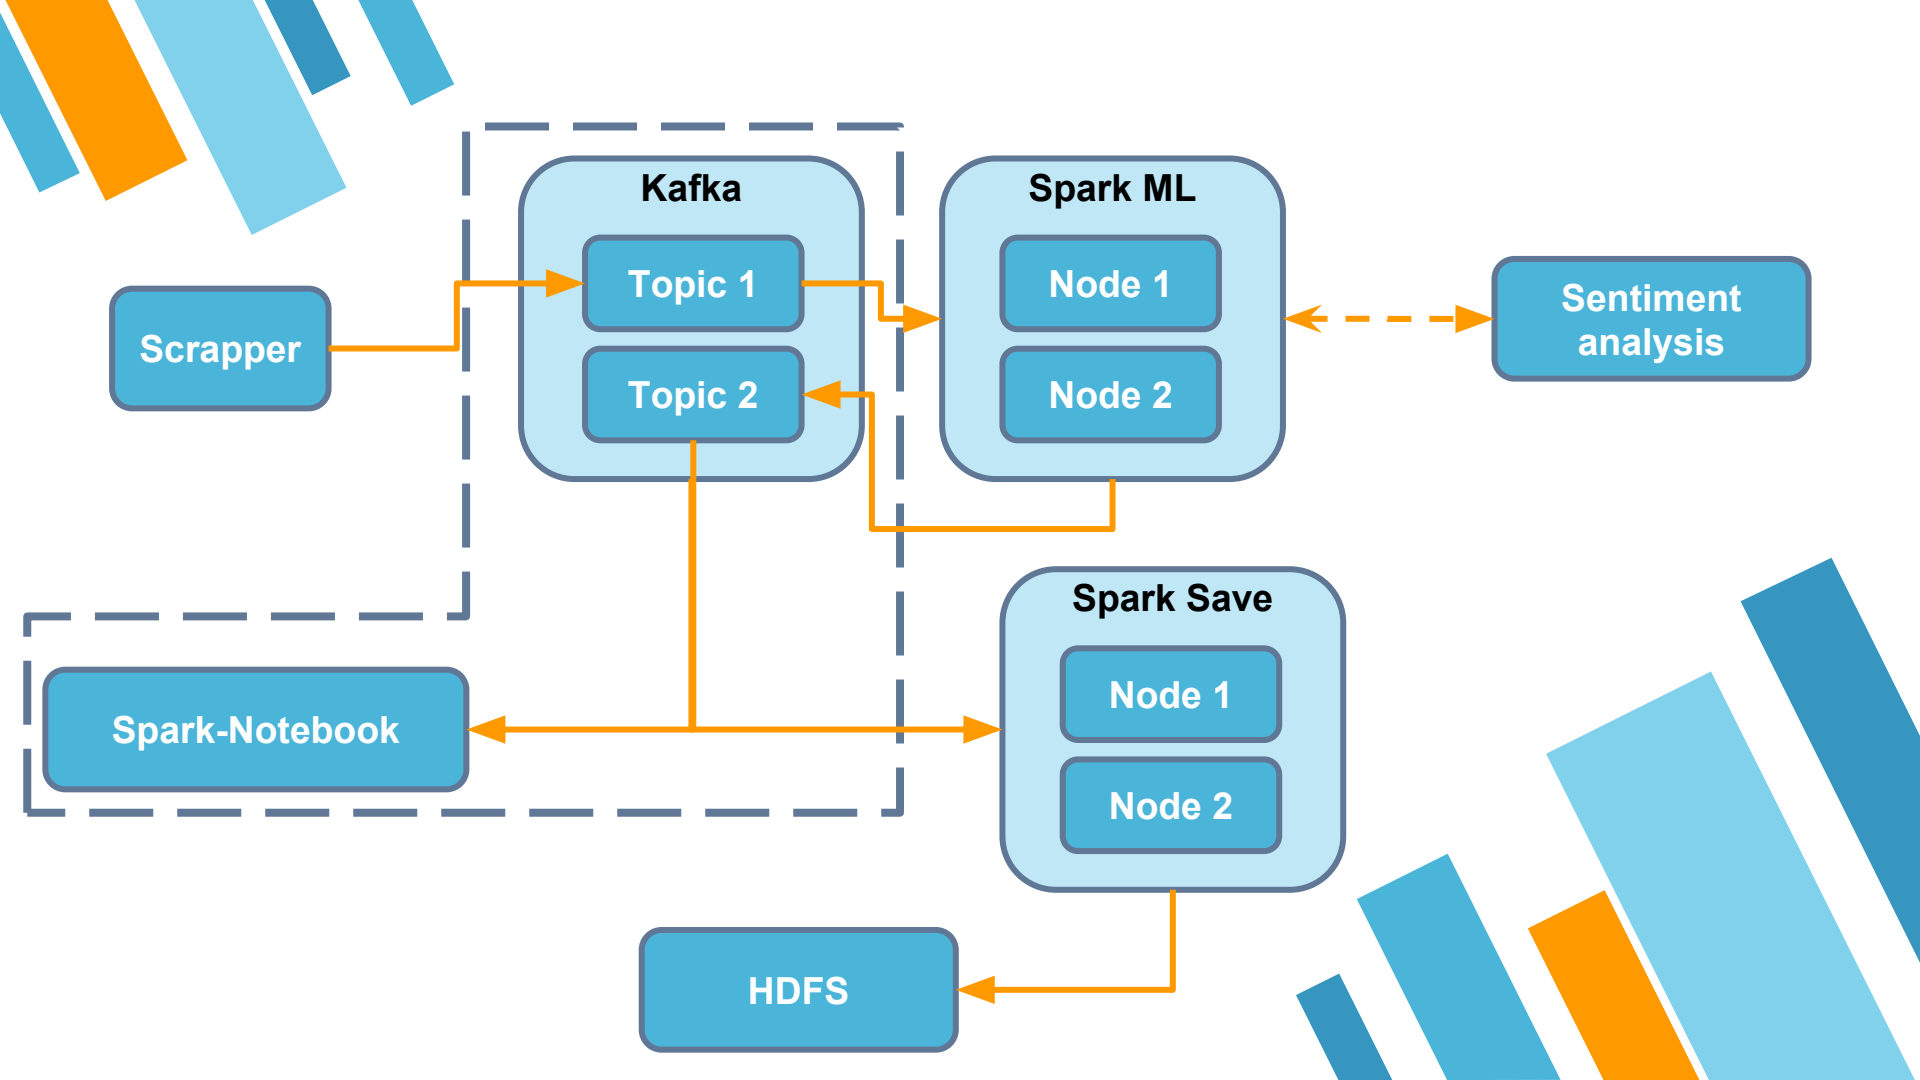
Kafka
Topic 1
Topic 2
Spark ML
Node 1
Node 2
Scrapper
Sentiment analysis
Spark Save
Node 1
Node 2
Spark-Notebook
HDFS

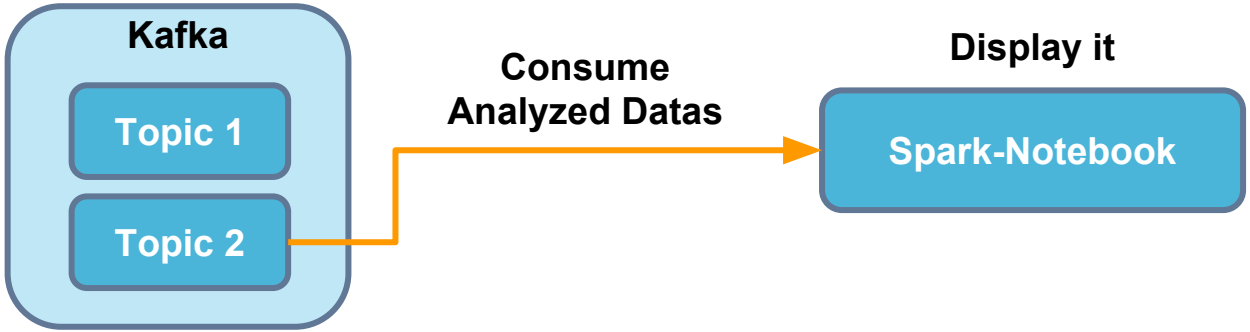
Kafka
Topic 1
Topic 2
Consume
Analyzed Datas
Display it
Spark-Notebook

SPARK-NOTEBOOK

- Easy to the data from kafka with spark streaming
- Prints cool graphics !

THANKS!

Any questions?